

# Lecture 9

# Demand Paging

Prof. Yinqian Zhang

Fall 2025

# Demand Paging Mechanisms

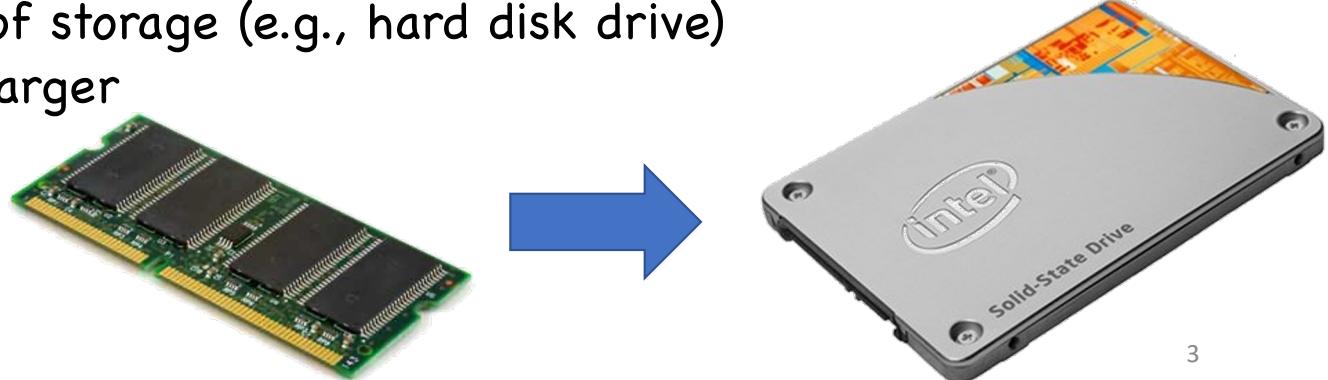
# How To Go Beyond Physical Memory?

- How to support large address space?
  - 64-bit machine supports up to 4EB address space
  - Applications may use more space than available in physical memory

现代 64 位机器支持极大的地址空间（理论上可达 4EB），但物理内存（RAM）通常只有 16GB 或 32GB

- Solution: stash away portions of address spaces that aren't currently in use
  - in the next-level of storage (e.g., hard disk drive)
  - slower but much larger

将当前不使用的地址空间部分“暂存”到下一级存储设备（如硬盘 HDD 或固态硬盘 SSD）中。虽然磁盘比内存慢很多，但容量大且便宜。



# An Abstraction of Address Space

- Who is responsible for moving data?
- Application: **memory overlays**
  - Application in charge of moving data *between memory and disk*
  - e.g., calling a function needs to make sure the code is in memory!
- OS: **demand paging**
  - OS configures page table entries
  - Virtual page maps to physical memory or files in disk
  - Process sees an abstraction of address space
  - OS determines where the data is stored

非常麻烦且容易出错

OS 配置页表 (Page Table)，将虚拟页面映射到物理内存或磁盘文件。  
对进程来说，它只看到一个连续的、巨大的抽象地址空间，并不关心数据到底在哪。

# Swap Space

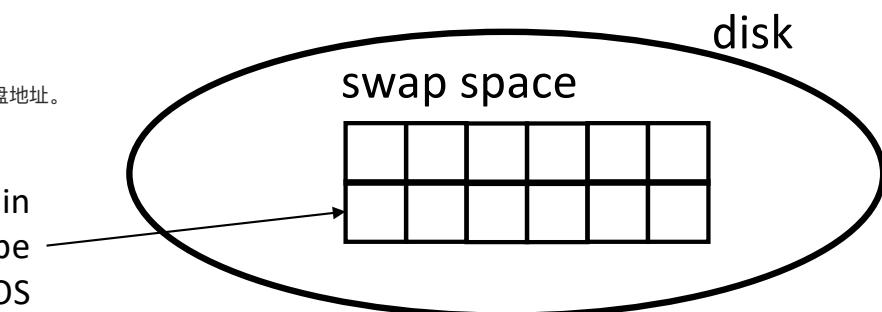
- Swap space is a partition or a file stored on the disk
  - OS swaps pages out of memory to it
  - OS swaps pages from it into memory
- Swap space conceptually divided into page-sized units
  - OS maintains a disk address of each page-sized unit

Swap Out ( 换出 ) : 当内存不足时 , OS 将不常用的页面从内存移动到 Swap Space。

Swap In ( 换入 ) : 当进程需要访问某个在 Swap 中的页面时 , OS 将其移回内存。

管理: Swap Space 在逻辑上也被切分成一个个“页大小”的单元 , OS 会记录每个单元的磁盘地址。

each page-sized unit in  
swap space can be  
addressed by the OS



# Swap Space Example

- 4-page physical memory and an 8-page swap space
  - Proc 0 has three virtual pages
  - Proc 1 has four virtual pages
  - Proc 2 and Proc 3 each has two virtual pages

	PFN 0	PFN 1	PFN 2	PFN 3				
Physical Memory	Proc 0 [VPN 0]	Proc 1 [VPN 2]	Proc 1 [VPN 3]	Proc 2 [VPN 0]				
	Block 0	Block 1	Block 2	Block 3	Block 4	Block 5	Block 6	Block 7
Swap Space	Proc 0 [VPN 1]	Proc 0 [VPN 2]	[Free]	Proc 1 [VPN 0]	Proc 1 [VPN 1]	Proc 3 [VPN 0]	Proc 2 [VPN 1]	Proc 3 [VPN 1]

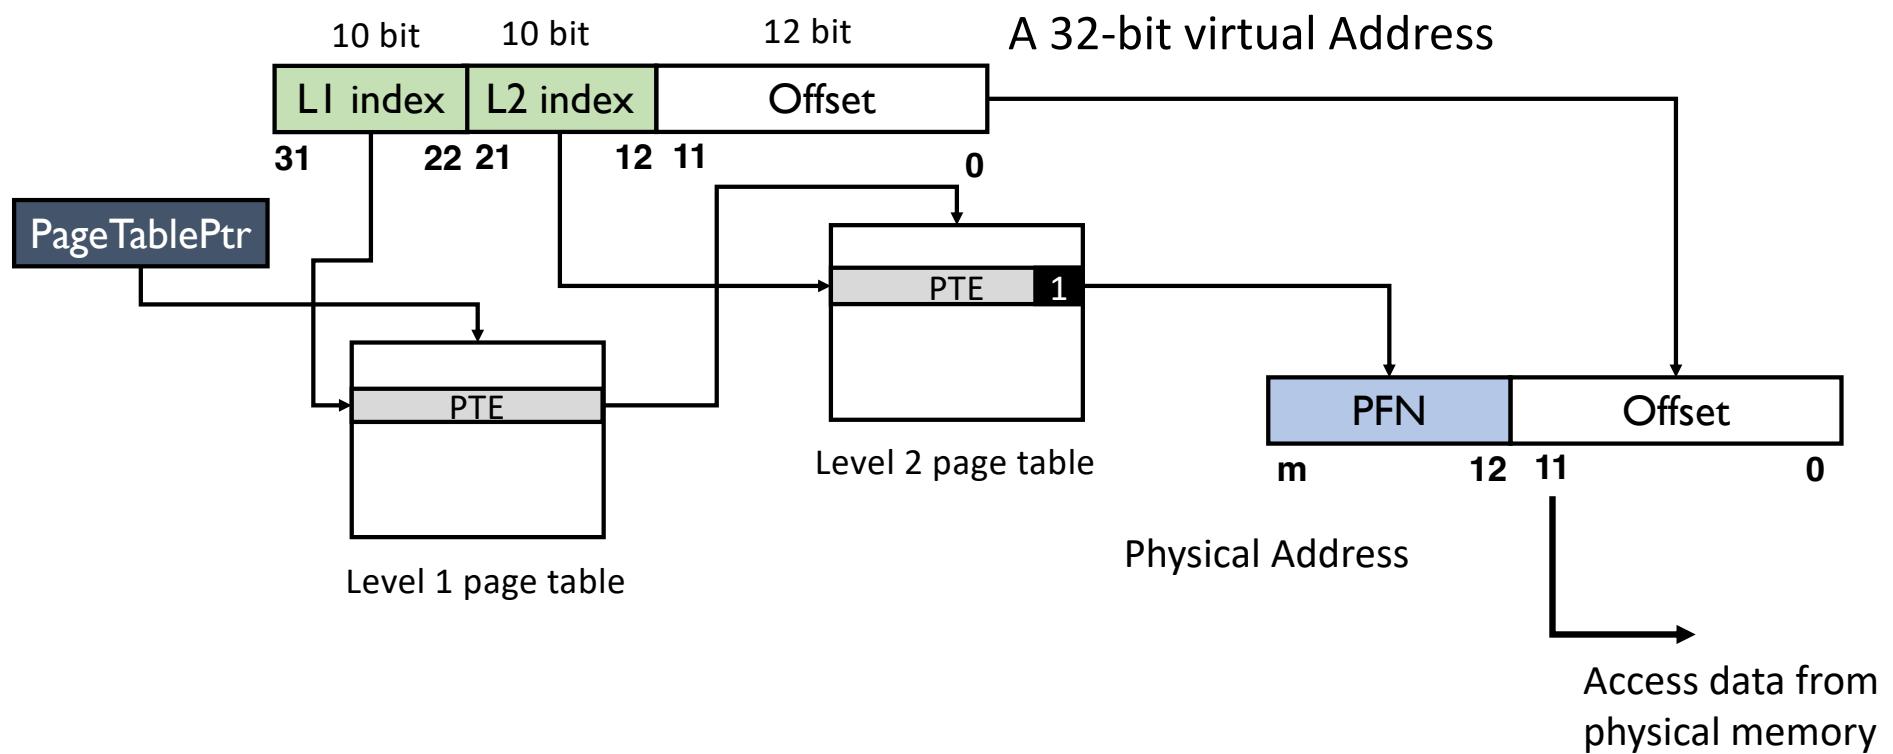
# Demand Paging

- Load pages from disk to memory **only as they are needed**
  - Pages are loaded “on demand”
- Data transferred in the unit of pages
- Two possible on-disk locations
  - Swap space:
    - created by OS for temporary storage of pages on disk
    - e.g., pages for stack and heap
  - Program binary files:
    - The code pages from this binary are only loaded into memory when they are executed
    - Read-only, thus never write back

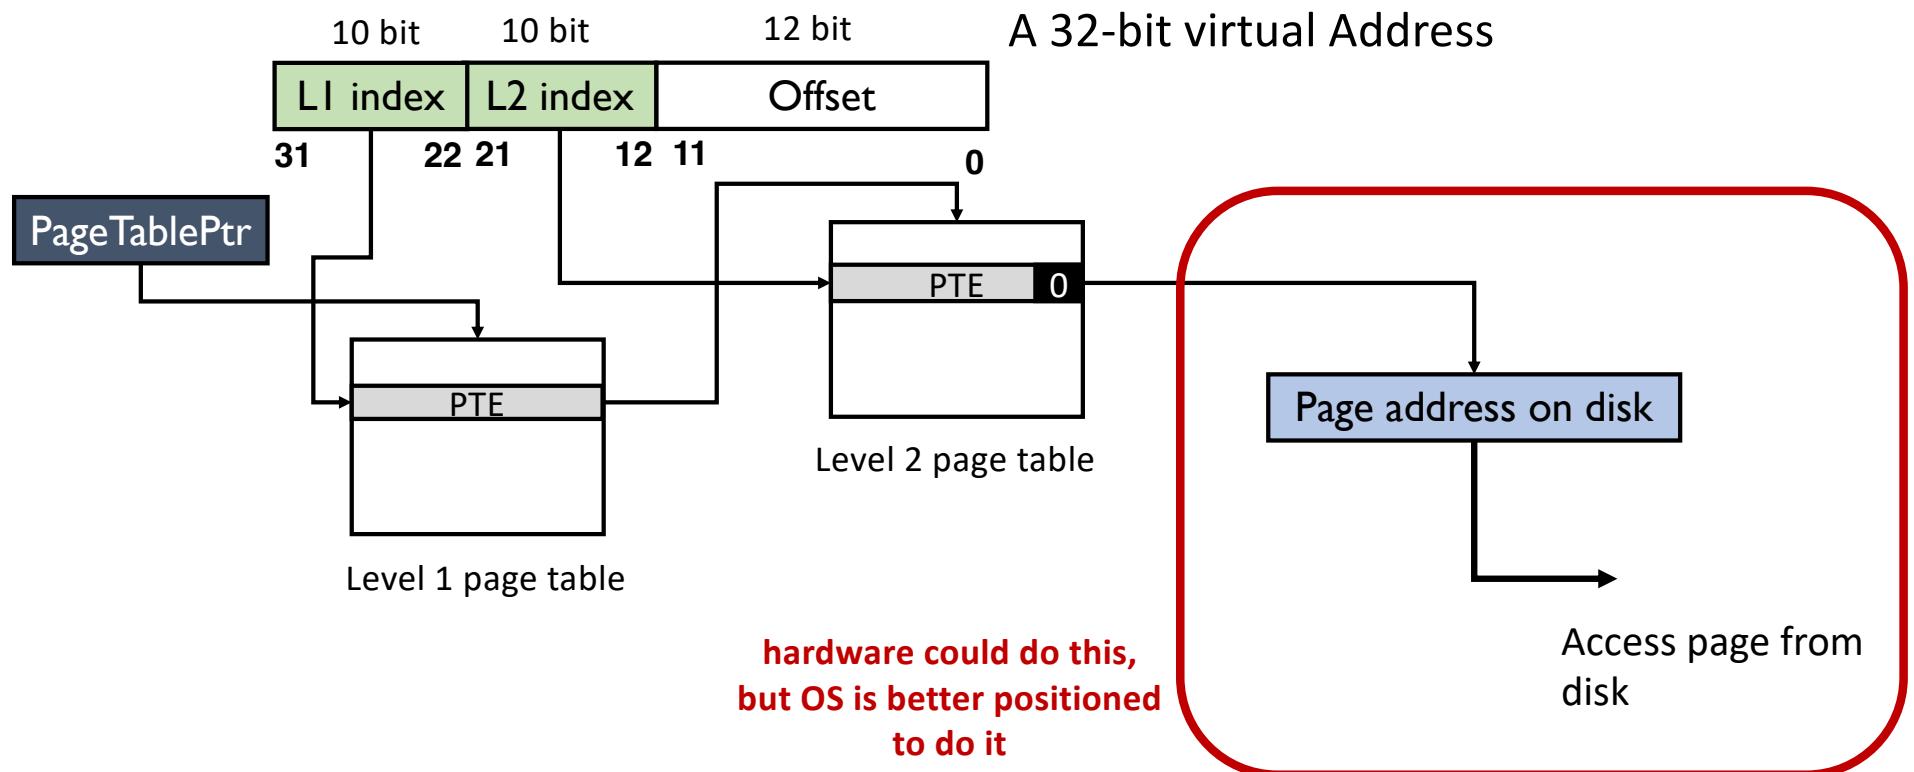
# Physical Memory as a Cache

- Physical memory can be regarded as a cache of on-disk swap space
- Block size of the cache?
  - 1 page (4KB)
- Cache organization (direct-mapped, set-associative, fully-associative)?
  - Fully associative: any disk page maps to any page frame
- What is page replacement policy?
  - LRU, Random, FIFO
- What happens on a miss?
  - Go to lower level to fill page (i.e. disk)
- What happens on a write, write-through or write back?
  - write-back: changes are written back to disk when page is evicted

# Present Bit



# Present Bit



操作系统读取这个磁盘地址，驱动磁盘控制器，将数据从磁盘（Swap Space 或文件系统）搬运到物理内存的一个空闲页中。10

# Page Faults

- Present bit = 0 raises a page fault exception
  - OS gets involved in address translation
- Page fault handler
  - (1) Find free page frame in physical memory
  - (2) Fetch page from disk and store it in physical memory
- After page fault
  - Return from page fault exception
  - CPU re-execute the instruction that accesses the virtual memory
  - No more page fault since present bit is set this time
  - TLB entry loaded from PTE

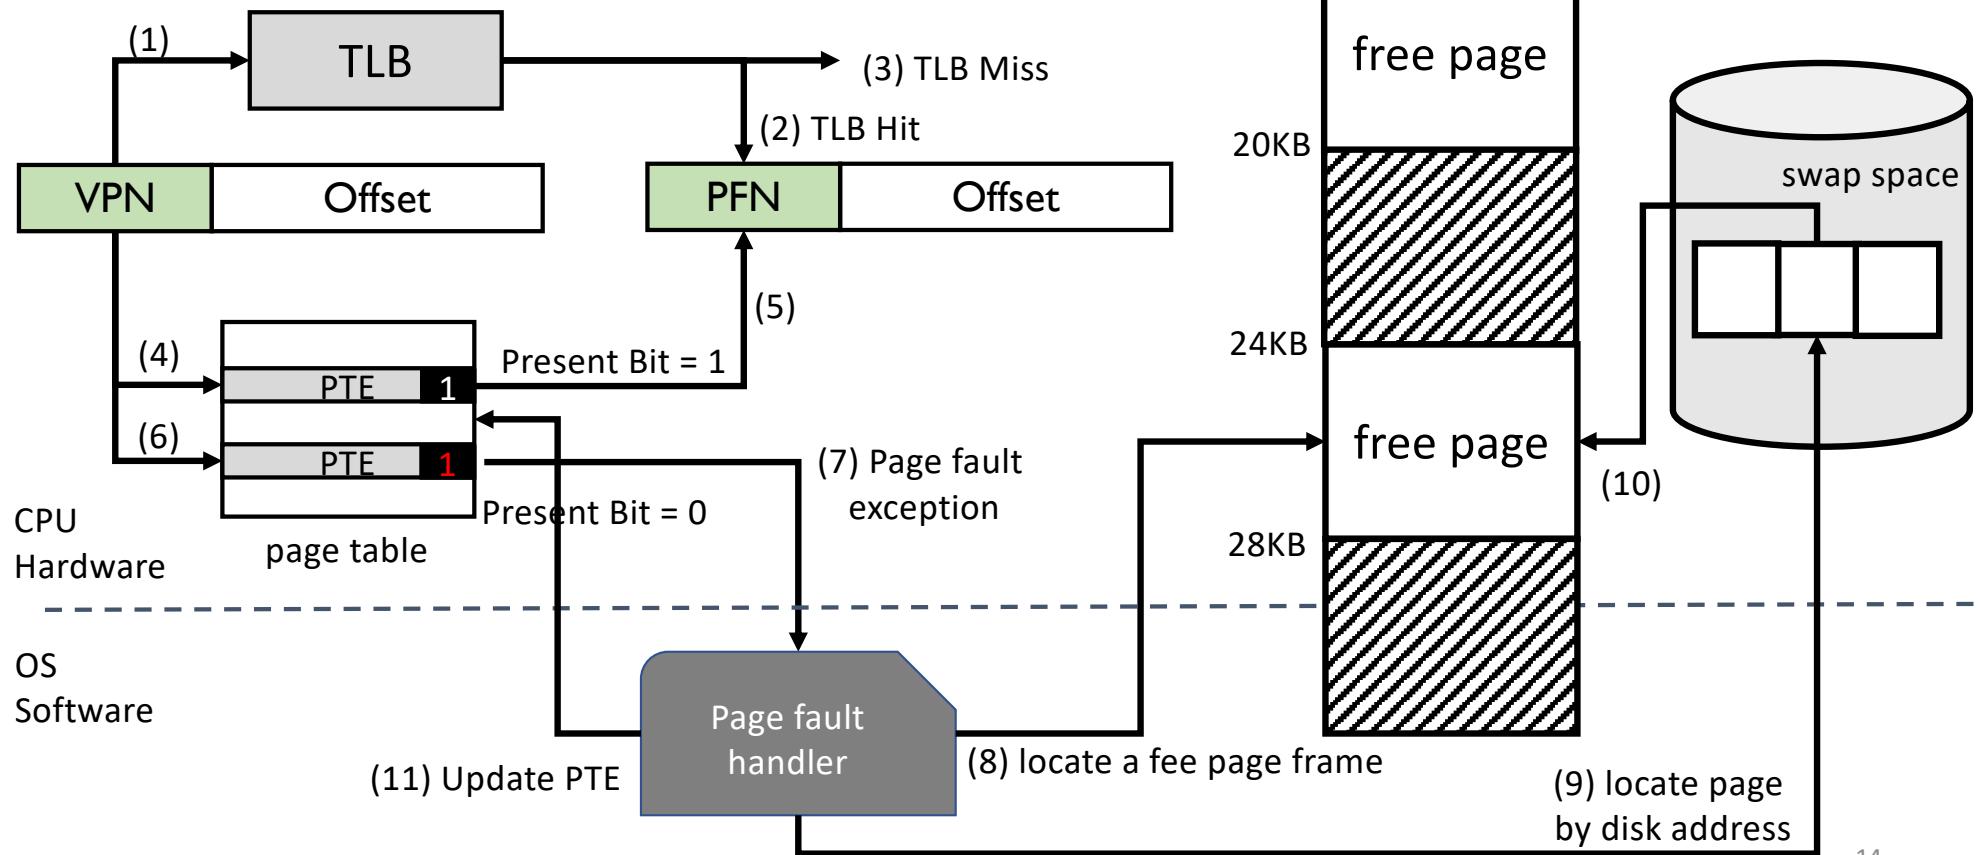
# Page Faults (Cont'd)

- (1) Find free page frame in physical memory
  - Find one free page frame from a free-page list
  - If no free page, trigger **page replacement** 内存页都满了, 需要找到一个内存页清理出去
- **Page Replacement**
  - find a page frame to be replaced
    - Page replacement policy decides which one to replace
  - If page frame to be replaced is **dirty**, 之前被修改过 write it back to disk
  - Update all PTEs pointing to the page frame
  - Invalidate all TLB entries for these PTEs

# Page Faults (Cont'd)

- (2) Fetch page from disk
  - Determine the faulting virtual address from register
  - Locate the disk address of the page in PTE (where PFN should be stored)
    - It is a very natural choice to make use of the space in PTE
  - Issues a request to disk to fetch the page into memory
  - Wait ..... (could be a very long time, context switch!)
  - When I/O completes, update page table entry: PFN, present bit

# Put It All Together



# When to Trigger Page Replacement

- **Proactive** page replacement usually leads to better performance

如果每次都要等到内存完全被占满（0 空闲）时才开始置换，那么那个不幸触发缺页的进程会非常痛苦。因为它不仅要等待数据从磁盘读进来（Swap In），还得先等待旧数据写回磁盘（Swap Out）。这会有双倍的延迟。

- Page replacement even though no one needs free page frames (yet)
- Always reserve some free page frames in the system

- Swap daemon

- background process for reclaiming page frames
- Low watermark: a threshold to trigger swap daemon
- High watermark: a threshold to stop reclaiming page frames

当 CPU 比较空闲或者内存紧张时，它就会醒来，检查内存，把那些很久没用的页（LRU）写回到磁盘（Swap Out），然后把物理内存标记为空闲。

# Page Replacement Policy

# Effective Access Time

- $EAT = \text{Hit Rate} \times \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$
- Example:
  - Memory access time = 200 nanoseconds
  - Average page-fault service time = 8 milliseconds
  - Suppose  $p$  = Probability of miss,  $1-p$  = Probability of hit
  - Then, we can compute EAT as follows:

$$\begin{aligned} EAT &= (1-p) \times 200\text{ns} + p \times 8 \text{ ms} \\ &= (1-p) \times 200\text{ns} + p \times 8,000,000\text{ns} \end{aligned}$$

# Effective Access Time (Cont'd)

$$\begin{aligned} \text{EAT} &= (1-p) \times 200\text{ns} + p \times 8 \text{ ms} \\ &= (1-p) \times 200\text{ns} + p \times 8,000,000\text{ns} \end{aligned}$$

- If one access out of 1,000 causes a page fault, then EAT is about  $8.2 \mu\text{s}$ :
  - This is a slowdown by a factor of 40!
- What if we want slowdown by less than 10%?
  - $200\text{ns} \times 1.1 < \text{EAT} \Rightarrow p < 2.5 \times 10^{-6}$
  - This is about 1 page fault in 400,000!

# Types of Cache Misses: Three Cs

- Compulsory Misses:
  - Cold-start miss: pages that have never been fetched into memory before
  - Prefetching: loading them into memory before needed
- Capacity Misses:
  - Not enough memory: must somehow increase available memory size
  - One option: Increase amount of DRAM (not quick fix!)
  - Another option: If multiple processes in memory: adjust percentage of memory allocated to each one!
- Conflict Misses:
  - fully-associative cache (OS page cache) does not have conflict misses

# Page Replacement Policies

- Optimal (also called MIN):
  - Replace page that will not be used for the longest time
  - Lead to minimum page faults in theory
- FIFO (First In, First Out)
  - Throw out oldest page first
  - May throw out heavily used pages instead of infrequently used
- RANDOM:
  - Pick random page for every replacement
  - Pretty unpredictable – makes it hard to make real-time guarantees

# Replacement Policies (Con't)

- Least Recently Used (LRU):
  - Replace page that has not been used for the longest time
  - **Temporal locality of program**
    - If a page has not been used for a while, it is unlikely to be used in the near future
- Least Frequently Used (LFU)
  - Replace page that has not been accessed many times
  - **Spatial locality of program**
    - if a page has been accessed many times, perhaps it should not be replaced as it clearly has some value.

# Example: Optimal (MIN)

- Suppose we have 3 page frames, 4 virtual pages, and following reference string:
  - A B C A B D A D B C B

Ref: Page:	A	B	C	A	B	D	A	D	B	C	B
1	A									C	
2		B									
3			C			D					

- MIN: 5 faults
  - Where will D be brought in? Look for page not referenced farthest in future

# Example: FIFO

- Suppose we have 3 page frames, 4 virtual pages, and following reference string:
  - A B C A B D A D B C B

Ref:	A	B	C	A	B	D	A	D	B	C	B
Page:											
1	A					D				C	
2		B					A				
3			C						B		

- FIFO: 7 faults
- When referencing D, replacing A is bad choice, since need A again right away

# Example: LRU

- Suppose we have 3 page frames, 4 virtual pages, and following reference string:
  - A B C A B D A D B C B

Ref: Page:	A	B	C	A	B	D	A	D	B	C	B
1	A									C	
2		B									
3			C			D					

- LRU performs the same as Optimal

# Is LRU Always Close to Optimal?

- Consider the following reference string: A B C D A B C D A B C D
- LRU performs as follows (the same as FIFO):

Ref:	A	B	C	D	A	B	C	D	A	B	C	D
Page:												
1	A			D			C			B		
2		B			A			D			C	
3			C			B			A			D

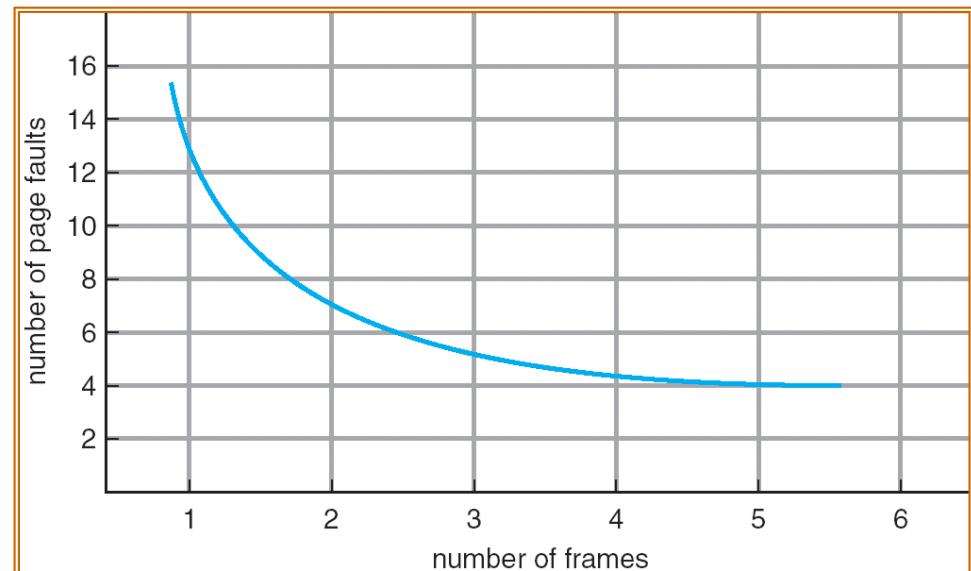
# Is LRU Always Close to Optimal? (Cont'd)

- Consider the following: A B C D A B C D A B C D
- MIN performs better:

Ref:	A	B	C	D	A	B	C	D	A	B	C	D
Page:												
1	A									B		
2		B						C				
3			C	D								

# Bélády's Anomaly

- One desirable property:  
When you add memory the miss rate drops
  - Yes for LRU and MIN
  - Not necessarily for FIFO!
- Bélády's anomaly
  - For FIFO, more page frames may lead to more page faults!



# Bélády's Anomaly Example

- Page replacement with 3 page frames

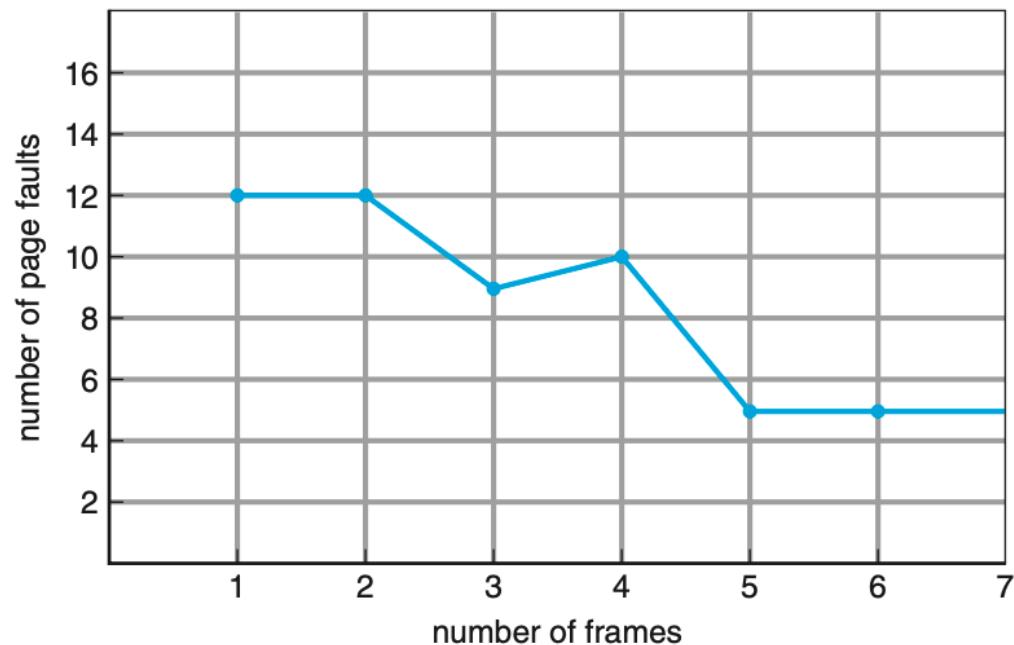
Ref: Page:	A	B	C	D	A	B	E	A	B	C	D	E
1	A			D			E					
2		B			A					C		
3			C			B					D	

- Page replacement with 4 page frames

Ref: Page:	A	B	C	D	A	B	E	A	B	C	D	E
1	A						E				D	
2		B						A				E
3			C						B			
4				D						C		

# Page Fault Curve

- Page fault curve for FIFO on reference string
  - 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



How do you plot a chart like this?

# LRU Implementation

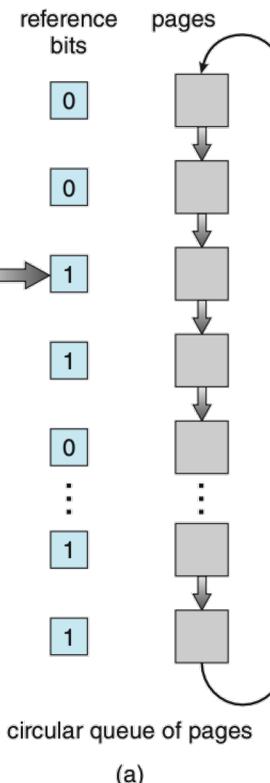
- Hardware support is necessary
  - Update a data structure in OS upon every memory access
  - E.g., a timestamp counter for each page frame
- Overhead
  - One additional memory write for each memory access
    - TLB hit does not save the extra memory access
  - Scan the entire memory to find the LRU one
    - 4GB physical memory has 1 million page frames
    - sorting is time consuming

# LRU Approximation with Reference Bit

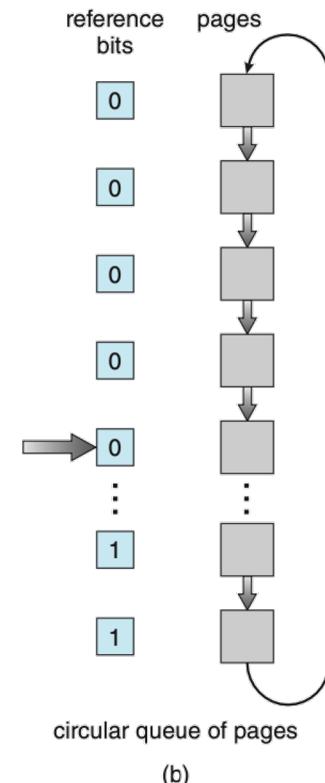
- Reference bit
  - One reference bit per page frame
  - All bits are cleared to 0 initially
  - The first time a page is referenced, the reference bit is set **by CPU**
    - Can be integrated with page table walk
  - The order of page accesses approximated by two clusters: **used** and **unused** pages
- Examples:
  - Clock algorithm (also called second-chance algorithm)
  - Enhanced clock algorithm with dirty bits

# Clock Algorithm

- Arrange physical pages in a circular list
- CPU sets reference bit to 1 upon first access
- OS maintains a pointer
  - When a replacement occurs, check reference bit of the current page
  - If 1: the page has been accessed recently, clear the bit (set to 0) and move to the next page
  - If 0: the page has not been accessed recently, good candidate for replacement, stop



(a)



(b)

# Clock Algorithm with Dirty Bit

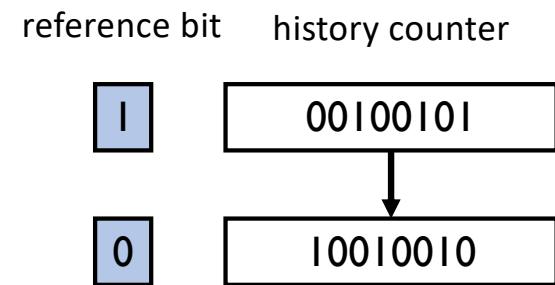
- Enhance clock algorithm with a dirty bit
  - dirty bit = 1: the page has recently been modified
- CPU sets dirty bit to 1 upon **write** access
- When a replacement occurs, OS checks (ref bit, dirty bit), and selects a candidate page in decreasing order
  - (0, 0) neither recently used nor modified — best page to replace
  - (0, 1) not recently used but modified — not quite as good, because the page will need to be written out before replacement
  - (1, 0) recently used but clean — probably will be used again soon
  - (1, 1) recently used and modified — probably will be used again soon, and the page will be need to be written out to secondary storage before it can be replaced

# LRU Approximation with Reference Bit and Counter

- Each physical page frame is associated with one reference bit and a counter
  - Reference bit indicate recent access
    - set by CPU hardware, cleared by OS
  - Counter records history of accesses
    - Maintained by OS
- Examples
  - Additional-reference-bits algorithm
  - $N^{\text{th}}$ -chance clock algorithm

# Additional-reference-bits Algorithm

- 8-bit history register associated with each page frame
- Timer interrupt every 100ms
  - reference bit shifts to highest bit in the history counter
  - other bits shift right and discard the lowest bit
  - 00000000 unused page in 800ms
- Compare history counter as unsigned integer
  - Larger value more recently used
  - $11000100 > 01110111$
- Approximate LRU with more bits and more frequent interrupts



# **N<sup>th</sup>-chance Clock Algorithm**

- All page frames arranged in a circular list and each page frame is associated with a reference bit and a counter
- CPU hardware sets reference bit upon memory accesses
- OS checks the reference bit of the page pointed to by the clock hand
  - 1 → clear reference bit and the counter
  - 0 → increment counter; if count=N, replace page
- How do we pick N?
  - Large N? Better approximation to LRU
    - If N ~ 1K, really good approximation
  - Small N? More efficient
    - Otherwise might have to look a long way to find free page

# Page Frame Allocation

# Allocation of Page Frames

- How do we allocate memory among different processes?
  - Does every process get the same fraction of memory? Different fractions?
  - Should we completely swap some processes out of memory?
- Minimum number of pages per process
  - Depends on the computer architecture
  - How many pages would one instruction use at one time
    - x86 only allows data movement between memory and register and no indirect reference
    - needs at least one instruction page, one data page, and some page table pages
- Maximum number of pages per process
  - Depends on available physical memory

# Global versus Local Allocation

- Global replacement
  - Process selects replacement frame from all page frames
  - One process can take a frame from another process
- Local replacement
  - Each process selects from only its own set of allocated frames

# Allocation Algorithms

- Equal allocation:
  - Every process gets same amount of memory
  - Example: 100 frames, 5 processes  $\rightarrow$  process gets 20 frames
- Proportional allocation
  - Number of page frames proportional to the size of process
  - $s_i$  = size of process  $p_i$  and  $m$  = total number of frame
  - $a_i$  = allocation for  $p_i$  =  $m \times \frac{s_i}{\sum s_j}$
- Priority Allocation:
  - Number of page frames proportional to the priority of process
  - Possible behavior: If process  $p_i$  generates a page fault, select for replacement a frame from a process with lower priority number

# Thrashing

- The memory demands of the set of running processes simply exceeds the available physical memory
- Early OS
  - Working set: the pages used actively of a process
  - Reduce the # of process so their working set fit into memory
- Modern OS
  - Out-of-memory killer when memory is oversubscribed
  - May need a reboot

# Thank you!

