

# RAID: Redundant Arrays of Inexpensive Disks

Team Members: Yuhan Rui, Yibin Lou , Youcheng Fang

December 26, 2025

# **Basics of RAID (from "Three Easy Pieces, CH38")**

# What is RAID?

- ▶ **Definition:** Redundant Array of Inexpensive Disks.
- ▶ **Origins:**
  - ▶ Introduced in the late 1980s.
  - ▶ By researchers at **U.C. Berkeley**: David Patterson, Randy Katz, and Garth Gibson.

## The Crux: How to make a Large, Fast, Reliable Disk?

- ▶ **Faster:** I/O operations are slow and can be the system bottleneck.
- ▶ **Larger:** Data grows rapidly; single disks get full quickly.
- ▶ **More Reliable:** If a disk fails without backup, valuable data is gone.
- ▶ **Solution:** Use multiple disks in concert to build a better storage system.

# RAID: Interface and Transparency

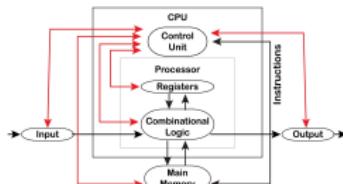
- ▶ **Externally:** Looks like a typical disk (a linear array of blocks).
- ▶ **Internally:** A complex system (a "computer" itself):
  - ▶ Multiple disks.
  - ▶ Memory (Volatile and Non-volatile).
  - ▶ One or more processors to manage the array.
- ▶ **Advantages:**
  - ▶ **Performance:** Parallel I/O speeds up operations.
  - ▶ **Capacity:** Aggregates multiple disks for large datasets.
  - ▶ **Reliability:** Uses **redundancy** to tolerate disk loss.

## Tip: Transparency Enables Deployment

- ▶ RAID replaces a disk **transparently**: No changes needed to the host computer, OS, or applications.
- ▶ This compatibility was key to RAID's massive success ("Deployability").

# Hardware RAID: A Specialized Computer

## Under the Hood: Components & Architecture



- ▶ **I/O Processor:** Runs firmware, offloads RAID logic (e.g., parity calc) from host CPU.
- ▶ **DRAM Cache:** GBs of memory buffer to accelerate R/W performance.
- ▶ **BBU (Battery Backup):** Solves “Consistency Update Problem” by preserving cache on power loss.

- ▶ **Host Interface:** Connects to motherboard (PCIe); presents “Single Large Disk” to OS.
- ▶ **XOR Engine:** Dedicated hardware circuit for extreme speed XOR parity calc.
- ▶ **Disk Interface:** Connects backend physical drives (SATA / SAS).

# Enterprise Storage: External Disk Arrays

Scaling Up: When a Card Isn't Enough



- ▶ **Independence:** Standalone hardware cabinet connected via Fiber/SAS.
- ▶ **Backplane & Hot-Swap:** Massive circuit board; allows replacing drives *without* powering down.
- ▶ **Ultimate Transparency:**
  - ▶ Internally: 100+ drives, complex RAID-6.
  - ▶ Externally: Server sees just **one cable and one giant disk.** (Storage Virtualization)

*“Essentially, it strips the ‘RAID Controller’ and ‘Batch of Disks’ from the host to create a standalone, specialized ‘Storage Server’.”*

# Fault Model

To understand RAID, we must first define the **Fail-Stop Fault Model**.

- ▶ **Two States Only:**
  - ▶ **Working:** All blocks can be read or written perfectly.
  - ▶ **Failed:** The disk is permanently lost.
- ▶ **Critical Assumption: Immediate Detection**
  - ▶ We assume the RAID controller can **immediately observe** when a disk has failed.
  - ▶ *Example:* The controller receives an error code or detects a timeout.
- ▶ **What we ignore (for now):**
  - ▶ *Silent Data Corruption:* Disk returns bad data without error.
  - ▶ *Latent Sector Errors:* Only part of the disk fails.

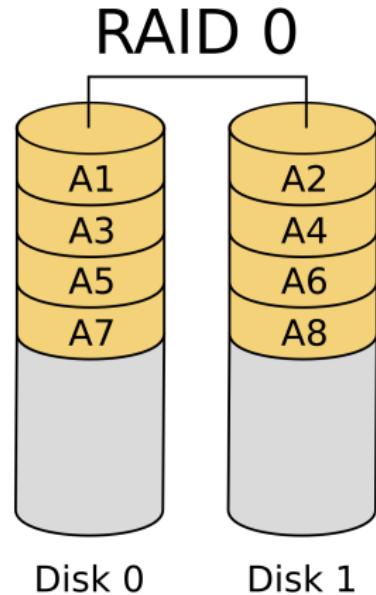
# Evaluation Metrics

## Three Axes of Analysis

- ▶ **1. Capacity:** How much *useful* space is available?
  - ▶ Given  $N$  disks with  $B$  blocks each.
  - ▶ **Range:** From  $N \cdot B$  (No redundancy) to  $N \cdot B/2$  (Mirroring).
- ▶ **2. Reliability:** How many disk faults can we tolerate?
  - ▶ Based on our **Fail-Stop** model (entire disk fails).
- ▶ **3. Performance:** The most challenging metric.
  - ▶ *Why?* It depends heavily on the **Workload**.
  - ▶ We must test against distinct patterns: Sequential vs. Random, Read vs. Write.

# RAID Level 0: Striping

Performance at all costs



- ▶ Not “True” RAID:
  - ▶ No redundancy at all.
  - ▶ Excellent **upper-bound** for performance and capacity.
- ▶ Round-Robin Distribution:
  - ▶ Spreads blocks across disks to extract max **parallelism**.
- ▶ Terminology:
  - ▶ Blocks in the same row (e.g., 0, 1, 2, 3) form a “**Stripe**”.

# RAID 0: Chunk Sizes & Mapping

## Fine-tuning the Striping Strategy

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |
|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      |
| 4      | 5      | 6      | 7      |
| 8      | 9      | 10     | 11     |
| 12     | 13     | 14     | 15     |

*Chunk Size = 1 Block (4KB)*

| Disk 0 | Disk 1 | Disk 2 | Disk 3 |                        |
|--------|--------|--------|--------|------------------------|
| 0      | 2      | 4      | 6      | chunk size:<br>1 block |
| 1      | 3      | 5      | 7      |                        |
| 8      | 10     | 12     | 14     |                        |
| 9      | 11     | 13     | 15     |                        |

*Chunk Size = 2 Blocks (8KB)*

## Impact of Chunk Size

- ▶ **Small:** High Parallelism.
- ▶ **Big:** High Concurrency.

## Mapping Formula (Generalized)

Let  $S$  = Chunk Size,  $N$  = Disks.

- ▶  $\text{ChunkIdx} = \lfloor A/S \rfloor$
- ▶  $\text{Disk} = \text{ChunkIdx \% } N$
- ▶  $\text{Offset} = \lfloor \text{ChunkIdx}/N \rfloor \cdot S + (A \% S)$

Example ( $S = 2$ ,  $N = 4$ , Block 14)

- ▶  $\text{ChunkIdx} = 14/2 = 7$
- ▶  $\text{Disk} = 7 \% 4 = 3$  (Disk 3)
- ▶  $\text{Offset} = \lfloor 7/4 \rfloor \cdot 2 + (14 \% 2) = 2$
- ▶ *Result: Disk 3, Offset 2*

# Chunk Size Trade-offs

## Parallelism vs. Positioning Time

- ▶ **Small Chunk Size** (e.g., 4KB):
  - ▶ **Pro:** Increases **Parallelism**. A single file stripes across many disks (fast transfer).
  - ▶ **Con:** Increases **Positioning Time**.
    - ▶ Accessing multiple disks means waiting for the *slowest* disk to position.
    - ▶  $T_{position} = \max(T_{disk1}, T_{disk2}, \dots)$
- ▶ **Big Chunk Size** (e.g., 64KB):
  - ▶ **Pro:** Reduces **Positioning Time**. A file fits on fewer disks (or just one).
  - ▶ **Con:** Reduces Intra-file Parallelism.
    - ▶ Throughput relies on **Concurrency** (multiple independent requests).
- ▶ **Conclusion:**
  - ▶ Determining the “best” size is hard; it requires knowing the **Workload**.
  - ▶ *Note: Real arrays often use 64KB, but we assume 4KB for simplicity.*

# RAID 0 Analysis & Performance Metrics

## Evaluating Striping and Defining Workloads

### RAID 0 Analysis

- ▶ **Capacity:** **Perfect** ( $N \cdot B$ ). No space wasted.
- ▶ **Reliability:** **Zero**. Any disk failure → Data Loss.
- ▶ **Performance:** **Excellent**. All disks utilize in parallel.
  
- ▶ **Performance Metrics:**
  - ▶ **Single-request Latency:** Time for one logical I/O operation.
  - ▶ **Steady-state Throughput:** Total bandwidth of concurrent requests (RAID is typically used in high-performance environments, so this is our main focus).

# Sequential vs. Random Workloads

Why  $S \gg R$ ?

## Disk I/O Time Components

For each request, I/O time can be decomposed into:

- ▶ **Seek Time:** Move the head to the right track.
- ▶ **Rotational Latency:** Wait for the sector to rotate under the head.
- ▶ **Transfer Time:** Actually move data between disk and memory.

## Workload Types

- ▶ **Sequential Access ( $S$  MB/s):**
  - ▶ Large, contiguous requests (e.g., scan a file from block  $x$  to  $x + 1\text{MB}$ ).
  - ▶ Disk stays on nearby tracks; little seeking/rotation, mostly transfer.
- ▶ **Random Access ( $R$  MB/s):**
  - ▶ Many small requests to unrelated locations (e.g., 4KB at block 10, then 550000, then 20100...).
  - ▶ Most time is spent in seek + rotation; little time in transfer.

In practice, sequential throughput  $S$  is **much larger** than random throughput  $R$  (i.e.,  $S \gg R$ ).

# Example: Computing $S$ and $R$

## A simple disk exercise

Assume disk parameters: seek 7 ms, rotation 3 ms, transfer rate 50 MB/s.

### Sequential throughput $S$

Request size: 10 MB

- ▶ Seek 7 ms, rotation 3 ms
- ▶ Transfer  $10 \text{ MB} / 50 \text{ MB/s} = 0.2 \text{ s} = 200 \text{ ms}$

$$\text{Total time} \approx 7 + 3 + 200 = 210 \text{ ms}$$

$$S = 10 \text{ MB} / 210 \text{ ms} \approx 47.6 \text{ MB/s.}$$

Most time is spent transferring data  $\Rightarrow S$  is near peak bandwidth.

### Random throughput $R$

Request size: 10 KB

- ▶ Seek 7 ms, rotation 3 ms
- ▶ Transfer  $10 \text{ KB} / 50 \text{ MB/s} \approx 0.195 \text{ ms}$

$$\text{Total time} \approx 7 + 3 + 0.195 \approx 10.195 \text{ ms}$$

$$R = 10 \text{ KB} / 10.195 \text{ ms} \approx 0.98 \text{ MB/s.}$$

Seek + rotation dominate  $\Rightarrow R < 1 \text{ MB/s, and } S/R \approx 50.$

**Takeaways:** Large sequential requests amortize seek/rotation and achieve high  $S$ ; tiny random requests are dominated by positioning time, so  $R$  is tiny and  $S \gg R$ .

# Back to RAID-0 Analysis

Putting  $S$  and  $R$  to use

## ► Latency (single request)

- ▶ For a single-block logical I/O, RAID-0 just sends the request to one physical disk.
- ▶ So latency is **about the same** as a single disk (seek + rotation + transfer on one disk).

## ► Steady-state throughput

- ▶ **Sequential ( $N \cdot S$  MB/s)**: use all  $N$  disks in parallel for large sequential requests.
- ▶ **Random ( $N \cdot R$  MB/s)**: with many independent random I/Os, all disks can still be busy.

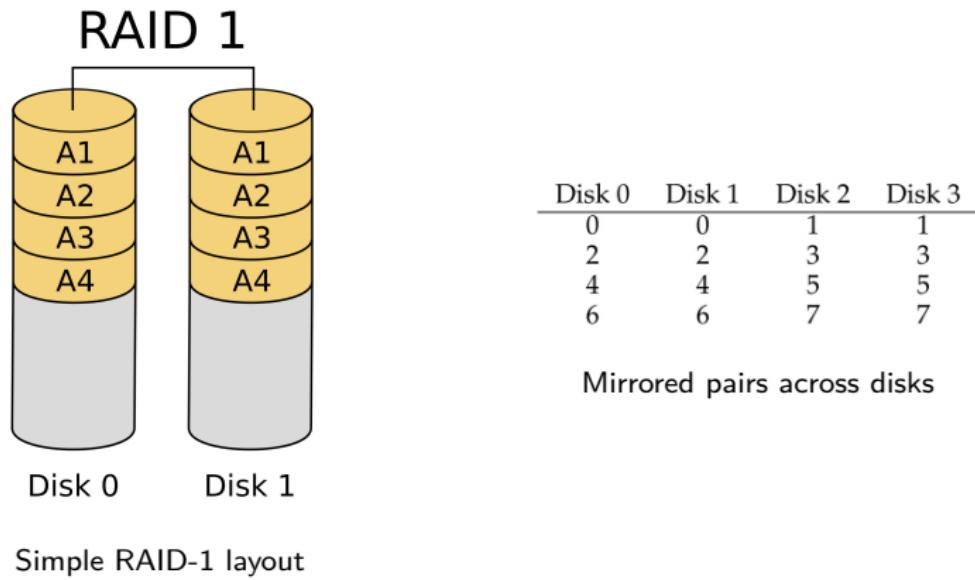
## ► Upper bound for comparison

- ▶ These simple formulas ( $N \cdot S$  and  $N \cdot R$ ) are easy to reason about.
- ▶ They serve as an **upper bound** when we compare to other RAID levels later.

# RAID Level 1: Mirroring

The expensive safety net

- ▶ **Layout:** Keep **two copies** of each logical block on **different disks**; disks form mirrored pairs (0–1, 2–3, ...), and data is **striped across pairs**.
- ▶ **Read/Write:** Reads can go to **either copy** of a block; writes must update **both** copies but can be issued in **parallel**.



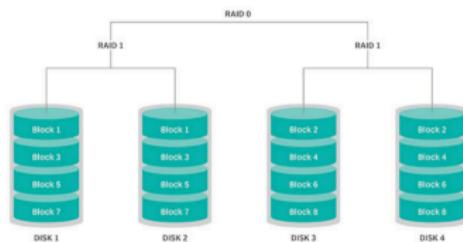
# RAID-10 vs. RAID-01

Stripe of mirrors vs. mirror of stripes

## RAID-10 (RAID 1+0)

RAID 10 (RAID 1+0)

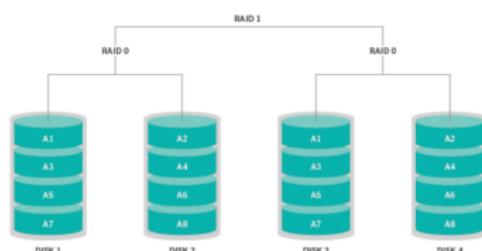
Stripe + Mirror



## RAID-01 (RAID 0+1)

RAID 01 (RAID 0+1)

Mirror + stripe



- ▶ Structure: combination of RAID 1 and RAID 0 → first **mirror** disks, then **stripe across mirrors**.
- ▶ **Pros:** RAID 0 performance + RAID 1 reliability.
- ▶ **Cons:** many disks (high cost), low space efficiency.

- ▶ Structure: combination of RAID 0 and RAID 1 → build RAID-0 stripes first, then **mirror the stripes**.
- ▶ **Pros:** RAID 0 performance + RAID 1 redundancy between stripes.
- ▶ **Cons:** **weaker** fault tolerance (any disk in a stripe kills that stripe) and low space efficiency.

# RAID 1 Analysis (I)

## Capacity, reliability, and latency

- ▶ **Capacity:** Mirroring level = 2  $\Rightarrow$  only half of raw space is useful.
  - ▶ With  $N$  disks of  $B$  blocks, useful capacity is  $\frac{N \cdot B}{2}$ .
- ▶ **Reliability:**
  - ▶ Can **tolerate 1 disk failure for sure** (per mirror pair).
  - ▶ In some lucky cases, can tolerate up to  $N/2$  failures (depending on which disks fail).
  - ▶ In practice, we assume it is “good for handling a single failure” and don’t rely on luck.
- ▶ **Single-request latency:**
  - ▶ **Read:** Same as a single disk. RAID-1 just directs the read to one copy.
  - ▶ **Write:** Needs two physical writes (to both copies) in parallel.
    - ▶ Worst-case seek/rotation of the two disks dominates  $\Rightarrow$  slightly worse than a single-disk write.

# RAID 1 Analysis (II)

## Steady-state throughput with $S$ and $R$

- ▶ **Sequential workloads:**
  - ▶ Each logical write  $\Rightarrow$  two physical writes (to both disks in a mirror pair).
  - ▶ Maximum sequential write bandwidth:  $\frac{N}{2} \cdot S$  MB/s (half of peak).
  - ▶ Sequential reads get the **same** bandwidth: still only  $\frac{N}{2} \cdot S$  MB/s.
    - ▶ Intuition: each disk only serves every other block (0,4,8,...)  $\Rightarrow$  each disk delivers only half of its peak.
- ▶ **Random workloads:**
  - ▶ Random reads: best case for RAID-1.
  - ▶ We can distribute requests across all  $N$  disks  $\Rightarrow$  throughput  $\approx N \cdot R$  MB/s.
  - ▶ Random writes: each logical write  $\Rightarrow$  two physical writes.
  - ▶ Effective random-write bandwidth  $\approx \frac{N}{2} \cdot R$  MB/s (half of what striping achieved).

## Aside: The RAID Consistent-Update Problem

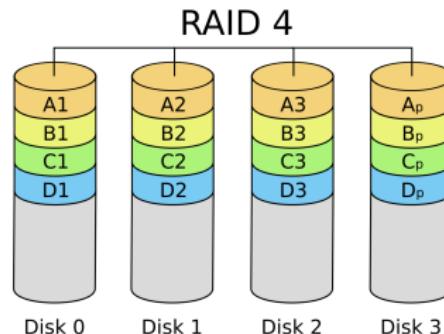
Making multi-disk writes atomic

- ▶ **Problem** (mirrored RAID): one logical write must update two disks. If power fails after Disk 0 is written but before Disk 1, the two copies become **inconsistent**.
- ▶ **Goal:** make multi-disk updates **atomic** – either both copies are updated, or neither is.
- ▶ **Solution:** use a **write-ahead log** in non-volatile storage.
  - ▶ Log what the RAID is about to do, then perform the disk writes.
  - ▶ After a crash, a **recovery** procedure replays or rolls back logged operations so mirrors end up consistent.

# RAID Level 4: Saving Space With Parity

Structure and basic idea

- ▶ **Goal:** Add redundancy like RAID-1, but use **less capacity**.
- ▶ **Design:** Data is striped across  $N - 1$  disks; the last disk stores a **parity block** for each stripe.
- ▶ **Parity idea:** For each stripe, compute 1 parity block  $P$  from data blocks using XOR.



| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      | P0     |
| 4      | 5      | 6      | 7      | P1     |
| 8      | 9      | 10     | 11     | P2     |
| 12     | 13     | 14     | 15     | P3     |

Example: one parity block per stripe  
 $(P0, P1, \dots)$

RAID-4 layout:  $N - 1$  data disks + 1 parity disk

# RAID Level 4: Parity Math

How parity and recovery work

- ▶ **Bit-level invariant:** For each row (data + parity), XOR of all bits is 0:  
 $C_0 \oplus C_1 \oplus C_2 \oplus C_3 \oplus P = 0$ .
- ▶ **Reconstruction:** If one disk/column (e.g.,  $C_2$ ) is lost, read the others +  $P$  and XOR them; the result is exactly the missing value (same XOR we used to create parity).
- ▶ **Blocks:** Real disks store 4KB blocks, so RAID just does this XOR bitwise across blocks: for each bit position, XOR data bits and write the result into the parity block.

| C0 | C1 | C2 | C3 | P                |
|----|----|----|----|------------------|
| 0  | 0  | 1  | 1  | XOR(0,0,1,1) = 0 |
| 0  | 1  | 0  | 0  | XOR(0,1,0,0) = 1 |

| Block0 | Block1 | Block2 | Block3 | Parity |
|--------|--------|--------|--------|--------|
| 00     | 10     | 11     | 10     | 11     |
| 10     | 01     | 00     | 01     | 10     |

# RAID 4 Analysis (I)

## Capacity and reliability

- ▶ **Capacity:** RAID-4 uses 1 disk for parity per RAID group.
  - ▶ With  $N$  disks of  $B$  blocks, useful capacity is  $(N - 1) \cdot B$ .
- ▶ **Reliability:**
  - ▶ Can tolerate **1 disk failure** in the group (any single disk).
  - ▶ If more than one disk in the group is lost, we cannot reconstruct the data.

| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      | P0     |
| 4      | 5      | 6      | 7      | P1     |
| 8      | 9      | 10     | 11     | P2     |
| 12     | 13     | 14     | 15     | P3     |

# RAID 4 Analysis (II)

Steady-state throughput: reads and full-stripe writes

- ▶ **Sequential reads:**
  - ▶ All  $N - 1$  data disks can stream data in parallel (parity disk mostly idle).
  - ▶ Peak sequential read bandwidth:  $(N - 1) \cdot S$  MB/s.
- ▶ **Sequential writes:** implemented as **full-stripe writes** (e.g., blocks 0,1,2,3 written together).
  - ▶ RAID computes new parity (e.g.,  $P_0$ ) as XOR of all new data blocks in the stripe.
  - ▶ Then writes all  $N$  blocks (data + parity) in parallel.
  - ▶ Effective sequential write bandwidth: also  $(N - 1) \cdot S$  MB/s.
- ▶ **Random reads:**
  - ▶ One-block random reads are spread across the  $N - 1$  data disks (not the parity disk).
  - ▶ Effective random-read throughput:  $(N - 1) \cdot R$  MB/s.

# RAID 4 Analysis (III)

## Random writes and parity update

- ▶ **Random writes:** overwrite a single data block in a stripe.
  - ▶ Must update both the data block and the parity block in that stripe.
- ▶ **Additive vs. subtractive parity** (bit view):
  - ▶ Additive: read all other data blocks in stripe, XOR with new data to form new parity (many reads when  $N$  is large).
  - ▶ **Subtractive:** use old data + new data + old parity:

$$P_{new} = (C_{old} \oplus C_{new}) \oplus P_{old}$$

- ▶ Same formula applies bitwise across an entire block (4KB, etc.).

# RAID 4 Analysis (IV)

Small-write I/O cost and parity bottleneck

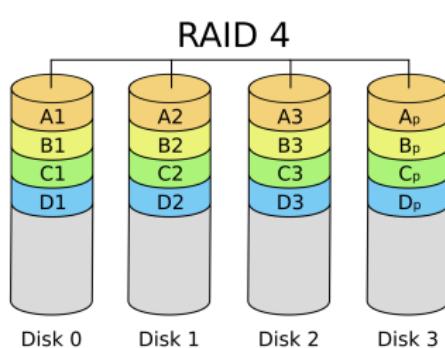
- ▶ I/O cost for small random writes (using subtractive parity):
  - ▶ Each logical write  $\Rightarrow$  2 reads (old data, old parity) + 2 writes (new data, new parity).
  - ▶ Parity disk does 2 I/Os per logical write (1 read, 1 write) and becomes a **bottleneck**.

| Disk 0 | Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------|--------|--------|--------|--------|
| 0      | 1      | 2      | 3      | P0     |
| *4     | 5      | 6      | 7      | +P1    |
| 8      | 9      | 10     | 11     | P2     |
| 12     | *13    | 14     | 15     | +P3    |

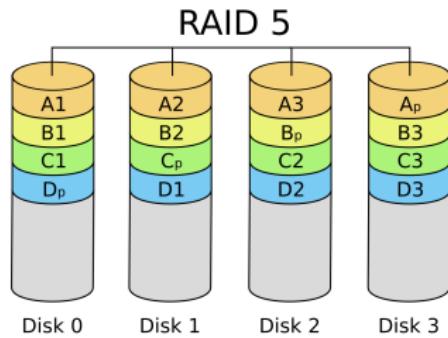
# RAID Level 5: Rotating Parity

From RAID-4 to RAID-5

- ▶ **Motivation:** RAID-4 small writes are limited by a single parity disk (bottleneck at  $R/2$  MB/s).
- ▶ **Key idea:** RAID-5 works like RAID-4, but **rotates the parity block across disks** for each stripe.
- ▶ **Effect:** spreads parity updates over all disks, removing the dedicated parity-disk bottleneck.



RAID-4: fixed parity disk



RAID-5: parity blocks rotated across disks

# RAID Level 5: Rotating Parity

## Fixing the RAID 4 Bottleneck

- ▶ **Solution:** Don't use a dedicated parity disk. **Rotate** the parity block assignment across all disks.
- ▶ **Impact:**
  - ▶ The “Parity Update” load is distributed evenly.
  - ▶ Random Write performance improves significantly because multiple stripes can be written in parallel.
- ▶ **Comparison:**
  - ▶ Same capacity as RAID 4:  $(N - 1) \times B$ .
  - ▶ Same reliability: Tolerates 1 failure.
  - ▶ **Much better random write concurrency.**

# RAID-5 Analysis

## How it compares to RAID-4

- ▶ **Capacity & reliability:** Same as RAID-4.
  - ▶ Effective capacity:  $(N - 1) \cdot B$ ; can tolerate 1 disk failure.
- ▶ **Sequential performance:**
  - ▶ Sequential read/write bandwidth and single-request latency are essentially the same as RAID-4.
- ▶ **Random reads:**
  - ▶ Slightly better than RAID-4, since all disks (including those holding parity for some stripes) can serve data for other stripes.
- ▶ **Random writes** (small writes):
  - ▶ Still need 4 I/Os per logical write (2 reads + 2 writes), but parity work is **spread across disks**.
  - ▶ With many independent writes, we can keep all  $N$  disks roughly busy.
  - ▶ Effective small-write bandwidth  $\approx \frac{N}{4} \cdot R$  MB/s – factor 1/4 comes from 4 I/Os per logical write.

# RAID Levels in Practice

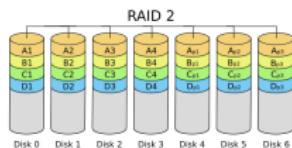
Choosing the “right” level

|                  | RAID-0      | RAID-1                                   | RAID-4                | RAID-5            |
|------------------|-------------|--|-----------------------|-------------------|
| Capacity         | $N \cdot B$ | $(N \cdot B)/2$                          | $(N - 1) \cdot B$     | $(N - 1) \cdot B$ |
| Reliability      | 0           | 1 (for sure)<br>$\frac{N}{2}$ (if lucky) | 1                     | 1                 |
| Throughput       |             |  |                       |                   |
| Sequential Read  | $N \cdot S$ | $(N/2) \cdot S^1$                        | $(N - 1) \cdot S$     | $(N - 1) \cdot S$ |
| Sequential Write | $N \cdot S$ | $(N/2) \cdot S^1$                        | $(N - 1) \cdot S$     | $(N - 1) \cdot S$ |
| Random Read      | $N \cdot R$ | $N \cdot R$                              | $(N - 1) \cdot R$     | $N \cdot R$       |
| Random Write     | $N \cdot R$ | $(N/2) \cdot R$                          | $\frac{1}{2} \cdot R$ | $\frac{N}{4} R$   |
| Latency          |             |  |                       |                   |
| Read             | $T$         | $T$                                      | $T$                   | $T$               |
| Write            | $T$         | $T$                                      | $2T$                  | $2T$              |

In practice, **different RAID levels** are chosen based on **workload, capacity, reliability, and cost trade-offs**, rather than a **single “best” option**.

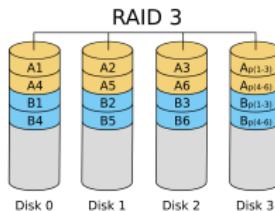
# Other RAID Levels

A quick structural tour



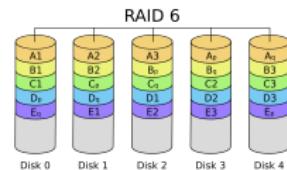
## RAID-2:

Uses bit-level striping with multiple  
**Hamming-code parity**  
disks;  
rarely used in practice.



## RAID-3:

Byte/bit-level striping with a  
**single shared parity** disk;  
good for large sequential  
I/O.



## RAID-6:

Like RAID-5 but with  
**two independent parity**  
blocks per stripe;  
tolerates 2 disk failures.

# Other Interesting RAID Issues

And a quick Part 1 wrap-up

- ▶ **Failure handling:** Hot spares, rebuild performance, and what happens while a disk is rebuilding.
- ▶ **More realistic faults:** Latent sector errors or block corruption, and techniques to detect/correct them.
- ▶ **Software RAID:** Implementing RAID in software can be cheaper but has its own issues (e.g., consistent-update).

**Part 1 takeaway:** RAID gives us powerful tools to trade off **capacity, reliability, and performance**, but picking the *right level and parameters* for a real workload is still **more of an art than a science**.

# **Part 2 — The Case for RAID (from “A Case for Redundant Arrays of Inexpensive Disks”)**

# Abstract

- ▶ **CPU & Memory Performance** increase.
- ▶ **I/O Performance:** If not matched, the performance gains of CPUs and memories will be wasted.
- ▶ **SLED (Single Large Expensive Disk):**
  - ▶ Capacity has grown rapidly.
  - ▶ Performance improvement has been modest.
- ▶ **RAID (Redundant Arrays of Inexpensive Disks):**
  - ▶ Based on magnetic disk technology developed for personal computers.
  - ▶ Offers an attractive alternative to SLED.
  - ▶ Promises improvements of an order of magnitude in:
    - ▶ Performance
    - ▶ Reliability
    - ▶ Power consumption
    - ▶ Scalability

# Growth of Computing and CPU Performance

- ▶ Gordon Bell's observation: 40% performance growth in single-chip computers per year (1974-1984).
- ▶ Bill Joy's prediction of even faster growth in 1985.
- ▶ The formula for MIPS growth:

$$\text{MIPS} = 2^{\text{Year}-1984}$$

- ▶ Amdahl: Each CPU instruction per second requires one byte of main memory

# Memory Growth

- ▶ Moore's Law: Transistor count doubles approximately every two years.
- ▶ Memory capacity has quadrupled every 2-3 years.
- ▶ The ratio of memory to MIPS is known as  $\alpha$ , with Amdahl's constant meaning  $\alpha = 1$ .
- ▶ Due to falling memory prices, the  $\alpha$  ratio has risen above 1 in modern systems.

# The Pending I/O Crisis

- ▶ Amdahl's Law:

$$S = \frac{1}{(1 - f) + \frac{f}{k}}$$

where:

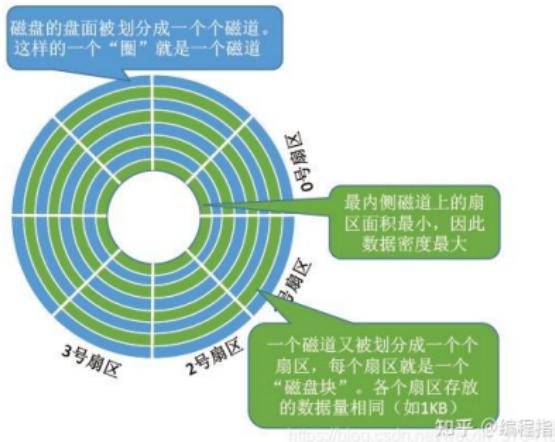
- ▶  $S$  = effective speedup,
- ▶  $f$  = fraction of work in faster mode,
- ▶  $k$  = speedup in the faster mode.

## Secondary Storage and Disk Technology

- ▶ The growth of secondary storage must match CPU and memory improvements.
- ▶ MAD (Maximal Areal Density): the maximum number of bits that can be stored per square inch, or the bits per inch in a track times the number of tracks per inch
- ▶  
$$\text{MAD} = 10^{(\text{Year}-1971)/10}$$
- ▶ Growth in disk capacity and corresponding price reduction.

# SLED Limitations

- ▶ The performance of SLED is dominated by **seek time** and **rotation latency**:
  - ▶ From 1971 to 1981, the raw seek time for a high-end IBM disk improved by only a factor of two.
  - ▶ Rotation latency remained unchanged.



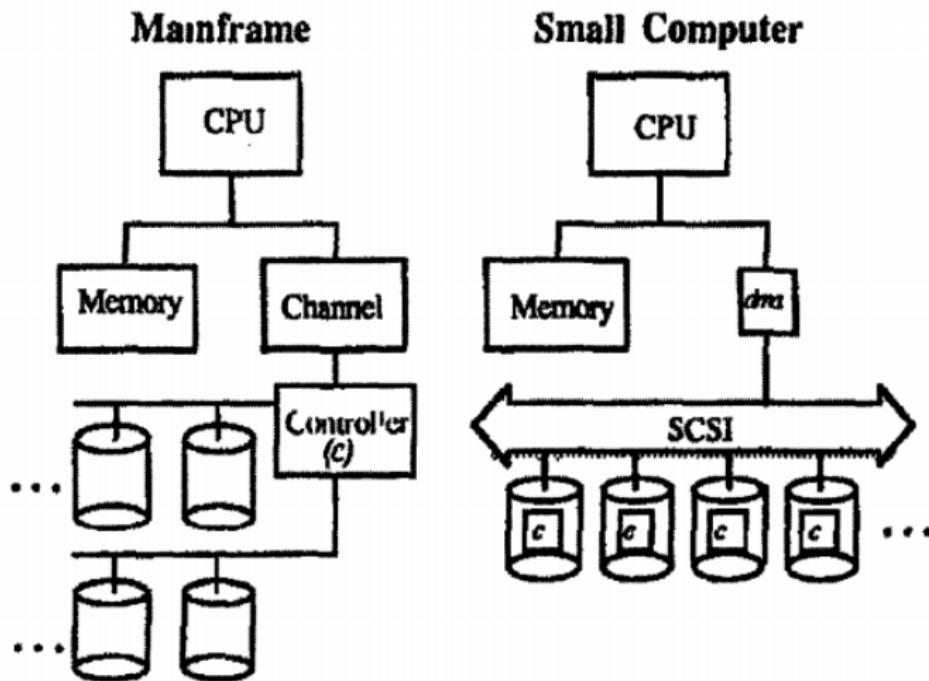
## A Solution: Arrays of Inexpensive Disks

Inexpensive disks achieve similar I/O performance per actuator as large disks.

| Characteristics               | IBM<br>3380 | Fujitsu<br>M2361A | Connors<br>CP3100 | 3380 v<br>3100 | 2361 v<br>3100 |     |
|-------------------------------|-------------|-------------------|-------------------|----------------|----------------|-----|
| Disk diameter (inches)        | 14          | 10.5              | 3.5               | 4              | 3              |     |
| Formatted Data Capacity (MB)  | 7500        | 600               | 100               |                | 01             | 2   |
| Price/MB(controller incl.)    | \$18-\$10   | \$20-\$17         | \$10-\$7          | 1-2.5          | 1              | 7-3 |
| MTTF Rated (hours)            | 30,000      | 20,000            | 30,000            |                | 1              | 1.5 |
| MTTF in practice (hours)      | 100,000     |                   | ?                 | ?              | ?              | ?   |
| No Actuators                  | 4           | 1                 | 1                 | 2              | 1              |     |
| Maximum I/O's/second/Actuator | 50          | 40                | 30                | 6              | 8              |     |
| Typical I/O's/second/Actuator | .30         | 24                | 20                | 7              | 8              |     |
| Maximum I/O's/second/box      | 200         | 40                | 30                | 2              | 8              |     |
| Typical I/O's/second/box      | 120         | 24                | 20                | 2              | 8              |     |
| Transfer Rate (MB/sec)        | 3           | 2.5               | 1                 | 3              | 4              |     |
| Power/box (W)                 | 6,600       | 640               | 10                | 660            | 64             |     |
| Volume (cu ft )               | 24          | 3.4               | 0.3               | 800            | 110            |     |

## A Solution: Arrays of Inexpensive Disks

SCSI make these disks highly cost-effective.



## The Bad News: Reliability

- ▶ MTTF (Mean Time To Failure) The average time a disk operates before failure.
- ▶ Assuming the failures of disks are independent and follow an exponential distribution.
- ▶ Since the failure rate of the array is the sum of the failure rates of individual disks:

$$\lambda_{\text{array}} = N \times \lambda_{\text{disk}}$$

- ▶ Since failure rate and MTTF are inversely proportional, the MTTF of the array is:

$$\text{MTTF of Array} = \frac{\text{MTTF of Single Disk}}{\text{Number of Disks in Array}}$$

## The Bad News: Reliability

### 1. Exponential CDF (Single Disk Failure Probability within t)

$$F(t) = P(X \leq t) = 1 - e^{-\lambda t}$$

### 2. Single Disk Reliability (Probability of Working within t)

$$P_{\text{single}}(t) = P(X > t) = e^{-\lambda t}$$

### 3. Total Reliability of N-Disk Array

$$P_{\text{array}}(t) = \prod_{i=1}^N P_{\text{single},i}(t) = e^{-N\lambda t}$$

### 4. Total Failure Rate of Array

$$\lambda_{\text{array}} = N\lambda$$

## RAID MTTF Calculation

$$\begin{aligned}MTTF_{RAID} &= \frac{MTTF_{Disk}}{G+C} * \frac{MTTF_{Disk}}{(G+C-1)*MTTR} * \frac{1}{n_G} \\&= \frac{(MTTF_{Disk})^2}{(G+C)*n_G * (G+C-1)*MTTR} \\MTTF_{RAID} &= \frac{(MTTF_{Disk})^2}{(D+C*n_G)*(G+C-1)*MTTR}\end{aligned}$$

# Reliability Overhead Cost and Usable Storage Capacity Percentage

**Reliability Overhead Cost:** This is the extra check disks, expressed as a percentage of the number of data disks  $D$ .

**Usable Storage Capacity Percentage:** The percentage of total capacity used for data storage.

# Performance in Supercomputers and Transaction-Processing Systems

## Supercomputers:

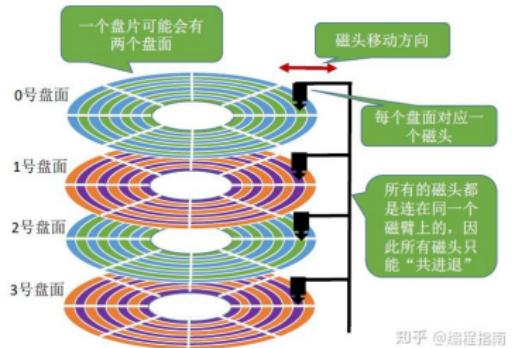
- ▶ The number of reads and writes per second for large blocks of data.

## Transaction-Processing Systems:

- ▶ The number of individual reads or writes per second.

## Summary:

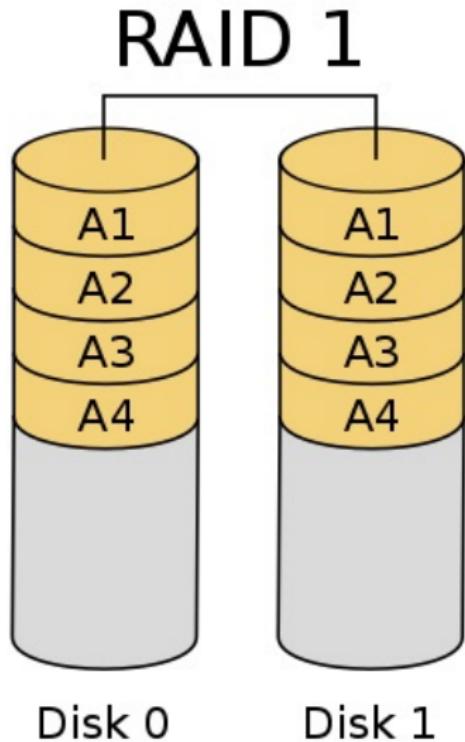
- ▶ Supercomputers: High data rate.
- ▶ Transaction-processing: High I/O rate.



知乎 @编程指南

## RAID 1: Mirrored Disks

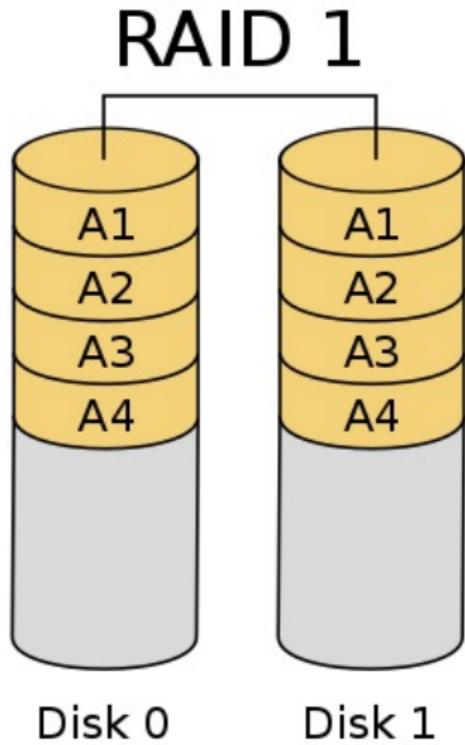
- ▶ Each data disk has an identical mirror, ensuring redundancy.
- ▶ Every write to a data disk is also written to the check disk.
- ▶ Provides high fault tolerance but requires double the disk capacity.



## RAID 1: Mirrored Disks

Tandem's optimization allows for:

- ▶ Parallel reads from both disks.
- ▶ Multiple controllers enabling optimized reading without waiting for the second write to complete.

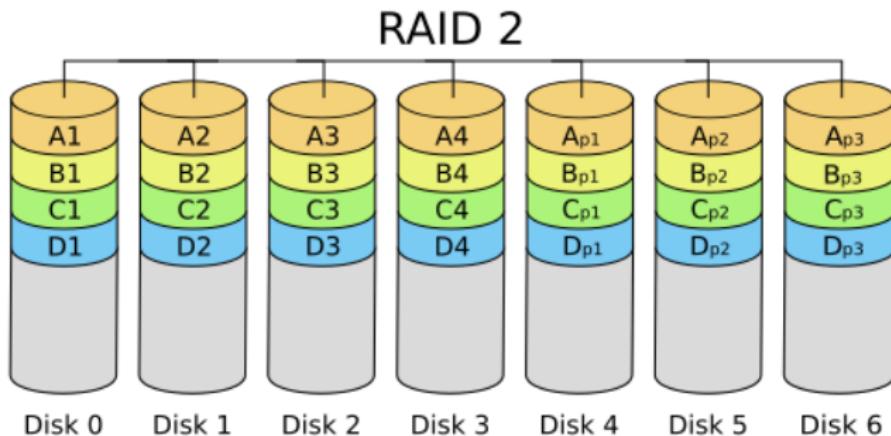


**Slowdown Factor:** A slowdown factor ( $S$ ) is introduced when more than two disks are involved in the group.

- ▶  $1 < S < 2$ , depending on the number of disks and the parallelism in use.

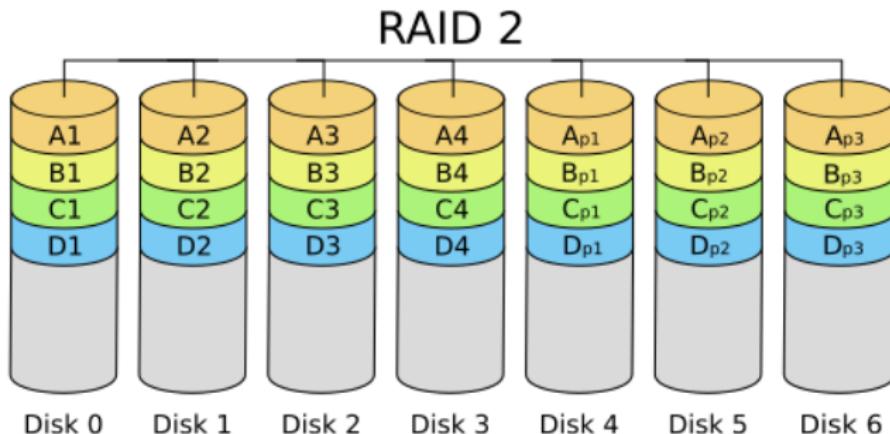
## RAID 2: Hamming Code for ECC

- ▶ It uses bit-interleaving and adds check disks to detect and correct errors.
- ▶ A single parity disk detects errors, but multiple check disks are required to correct errors.
- ▶ For 10 data disks (G), 4 check disks (C) are required; for 25 data disks, 5 check disks.



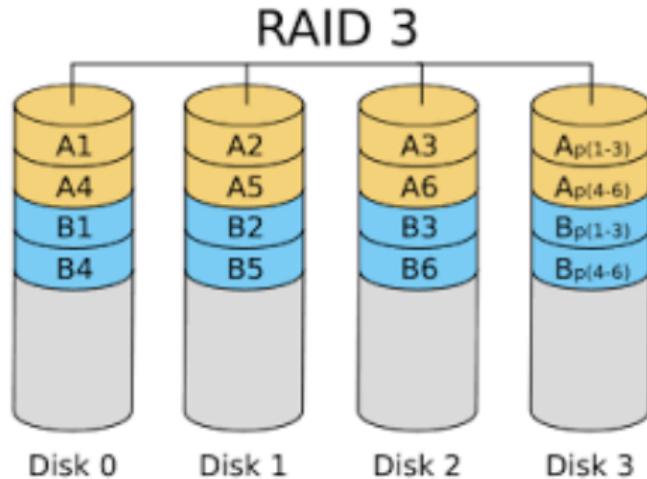
## RAID 2: Hamming Code for ECC

- ▶ Each data transfer unit is a sector.
- ▶ Bit-interleaved disks mean a large transfer requires at least G sectors.
- ▶ For small data transfers, reads and writes still imply reading the full sector from each bit-interleaved disk.
- ▶ Writes involve the read-modify-write (R-M-W) cycle across all disks in the group.



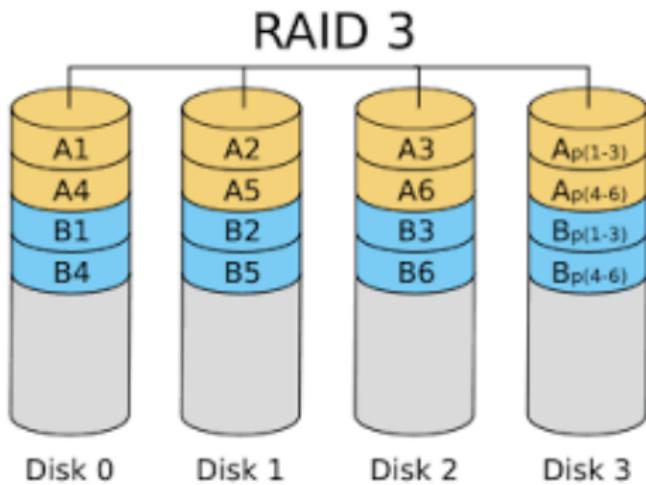
## Reducing Check Disks in RAID 3

- ▶ Reducing check disks to one per group ( $C = 1$ ).
- ▶ The performance of RAID 3 is the same as Level 2 RAID.
- ▶ The reduction in total disks increases reliability.



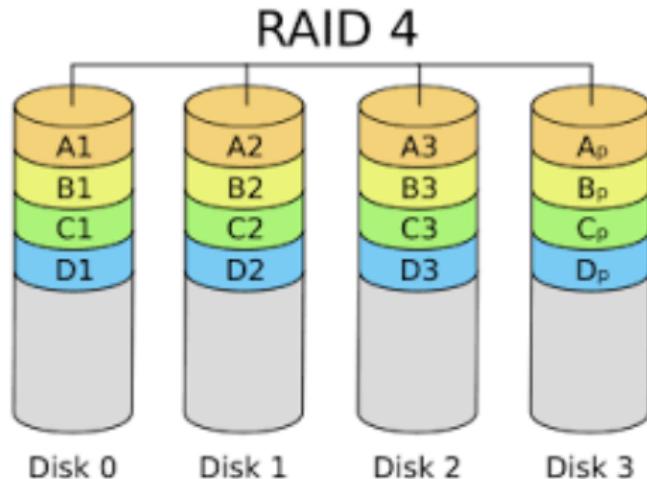
# Error Reconstruction

- ▶ **Error Detection:** Special signals provided in the disk interface or extra checking information at the end of a sector to detect and correct soft errors.
- ▶ **Error Correction:** Information on the failed disk can be reconstructed by calculating the parity of the remaining good disks.



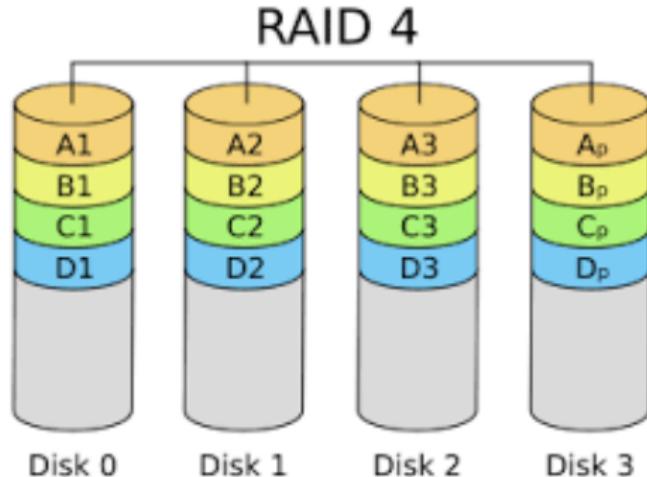
## RAID Level 4 - Independent Reads/Writes

- ▶ Unlike RAID 3, data is not spread across multiple disks at the bit level. Instead, data is stored in sectors, and each disk can handle its individual transfer.
- ▶ Reads can be performed independently on a single disk at maximum disk rate, while still detecting errors.



## RAID 4 Performance and Bottleneck

- ▶ Small writes improve since only 2 disks are involved in reading and writing the data and parity.
- ▶ The check disk is critical in limiting the number of simultaneous writes to the RAID system.

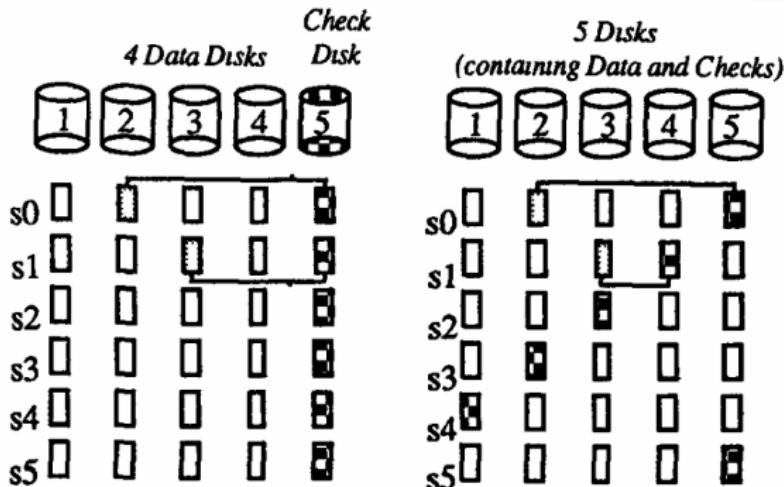


# Comparison of level 2,3,4

| 4 Transfer Units<br><i>a, b, c &amp; d</i> |  | a0    | b0                              | c0   | d0  |  |
|--|--|-------|---------------------------------|--|-----|--|
|  |  | a1    | b1                              | c1   | d1  |  |
|  |  | a2    | b2                              | c2   | d2  |  |
|  |  | a3    | b3                              | c3   | d3  |  |
| <i>Sector 0 Data Disk 1</i>                |  | a0    | a0                              | a0   | a0  |  |
|  |  | b0    | b0                              | b0   | b0  |  |
|  |  | c0    | c0                              | c0   | c0  |  |
|  |  | d0    | d0                              | d0   | d0  |  |
| <i>Sector 0 Data Disk 2</i>                |  | a1    | a1                              | b0   | b0  |  |
|  |  | b1    | b1                              | b1   | b1  |  |
|  |  | c1    | c1                              | c1   | c1  |  |
|  |  | d1    | d1                              | b2   | b2  |  |
| <i>Sector 0 Data Disk 3</i>                |  | a2    | a2                              | c0   | c0  |  |
|  |  | b2    | b2                              | b2   | b2  |  |
|  |  | c2    | c2                              | c2   | c2  |  |
|  |  | d2    | d2                              | c3   | c3  |  |
| <i>Sector 0 Data Disk 4</i>                |  | a3    | a3                              | d0   | d0  |  |
|  |  | b3    | b3                              | d1   | d1  |  |
|  |  | c3    | c3                              | d2   | d2  |  |
|  |  | d3    | d3                              | d3   | d3  |  |
| <i>Sector 0 Check Disk 5</i>               |  | aECC0 | ECCa                            | ECC0   |     |  |
|  |  | bECC0 | ECCb                            | ECC1   |     |  |
|  |  | cECC0 | ECCc                            | ECC2   |     |  |
|  |  | dECC0 | ECCd                            | ECC3   |     |  |
| <i>Sector 0 Check Disk 6</i>               |  | aECC1 | (Only one check disk in level 3 | (Each transfer unit is placed into a single sector                             |     |  |
|  |  | bECC1 |                                 | Note that the check info is now calculated over a piece of each transfer unit) |     |  |
|  |  | cECC1 |                                 |  | D   |  |
|  |  | dECC1 |                                 |  | I   |  |
| <i>Sector 0 Check Disk 7</i>               |  | aECC2 |                                 |  | S   |  |
|  |  | bECC2 |                                 |  | K   |  |
|  |  | cECC2 |                                 |  | (S) |  |
|  |  | dECC2 |                                 |  |     |  |

## RAID 5 Check Information Placement

- ▶ In RAID 4, check information is stored in a single check disk.
- ▶ In RAID 5, check information is distributed across all disks, including data disks.

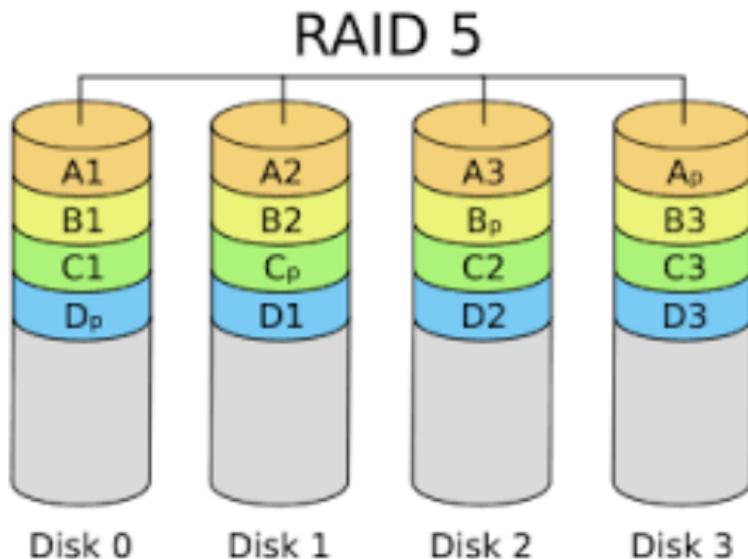


(a) Check information for Level 4 RAID for  $G=4$  and  $C=1$ . The sectors are shown below the disks. (The checked areas indicate the check information.) Writes

(b) Check information for Level 5 RAID for  $G=4$  and  $C=1$ . The sectors are shown below the disks, with the check information and data spread evenly through all the disks. Writes

## RAID 5 Parallel Write Advantage

- ▶ RAID 5 supports multiple independent writes per group.
- ▶ In RAID 4, writing to sectors on different disks requires sequential writes to the check disk.
- ▶ In RAID 5, writes to different sectors can occur in parallel, significantly improving write performance.



# RAID 5 vs SLED Comparison

- ▶ **Performance:** RAID 5 provides up to 10x better performance compared to IBM 3380, and 5x better than Fujitsu M2361A.
- ▶ **Reliability:** RAID 5 has significantly higher MTTF due to redundancy and data reconstruction.
- ▶ **Power Consumption and Size:** RAID 5 is more energy-efficient and smaller in size, reducing cooling costs.
- ▶ **Storage Capacity and Expandability:** RAID 5 offers better storage utilization and modular expansion compared to SLED.
- ▶ **Advantages:** RAID 5 is ideal for supercomputer applications, transaction processing, and systems with limited storage.

# Challenges: Data Corruption

(Based on: "An Analysis of Data Corruption in the Storage Stack", FAST '08)

# Data Integrity Segment (DIS)

Building the Foundation for End-to-End Protection

## 1. Data Structure:

- ▶ Every **4KB** data block is mandatorily attached with **64 bytes** of metadata.

## 2. Physical Implementation (Layout):

### Enterprise Class (FC)

- ▶ Uses **520-byte** sectors.
- ▶  $8 \times 520 = 4160$  bytes.
- ▶ *Result:* Perfect embedding, zero performance overhead.

### Nearline Class (SATA)

- ▶ Uses standard **512-byte** sectors.
- ▶ *Result:* Requires consuming extra sector space for DIS.

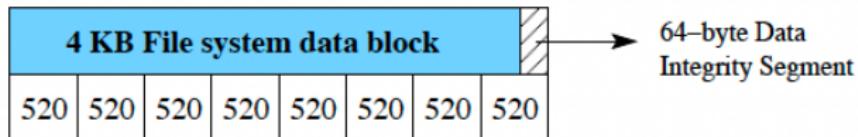
## 3. DIS Components:

- ▶ **Checksum:** Verifies physical bits.
- ▶ **Identity:** Verifies logical context (Inode, Offset).

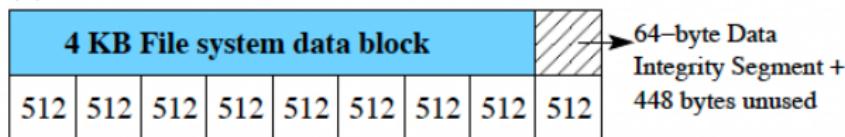
# Data Integrity Segment (DIS)

Building the Foundation for End-to-End Protection

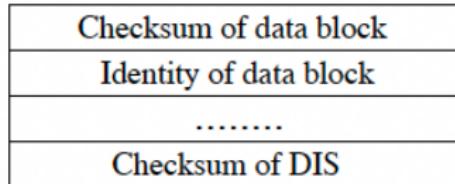
(a) Format for enterprise class disks



(b) Format for nearline disks



(c) Structure of the data integrity segment (DIS)



# Real-Time Defense: Detection on Read

## 1. Checksum Mismatch (CM)

- ▶ **Layer:** RAID / Adapter.
- ▶ **Logic:**  
 $\text{Hash}(\text{Data}) \neq \text{Stored}$ .
- ▶ **Target:** Physical media corruption, bit flips.

## 2. Identity Discrepancy (ID)

- ▶ **Layer:** File System.
- ▶ **Logic:** Checksum is OK, but DIS Inode  $\neq$  Request.
- ▶ **Target:** Lost Writes, Misdirected Writes.

### Recovery Action:

Failure in *either* check → Trigger **RAID Reconstruction**.

# Data Scrubbing

## Scrubbing Workflow:

1. **Verify Integrity:** Check Checksums of all blocks (finds latent CMs).
2. **Verify Logic:** Check RAID Parity Consistency.

## Unique Error: Parity Inconsistency (PI)

- ▶ **Phenomenon:** All Data Checksums are OK, but  $Data + Data \neq Parity$ .
- ▶ **Causes:** Memory corruption, Software bugs, non-atomic updates.
- ▶ **Resolution:** Rewrite Parity.

# Core Challenge: Silent Corruption

## Fail-Stop vs. Silent Corruption

### Traditional RAID (Fail-Stop)

- ▶ Disk Fails
- ▶ **Error Reported**
- ▶ RAID Rebuilds

### Reality (Silent Corruption)

- ▶ Disk Fails (Bit flip)
- ▶ **NO Error Reported**
- ▶ RAID propagates bad data

# Type 1: Checksum Mismatches (CMs)

## Definition

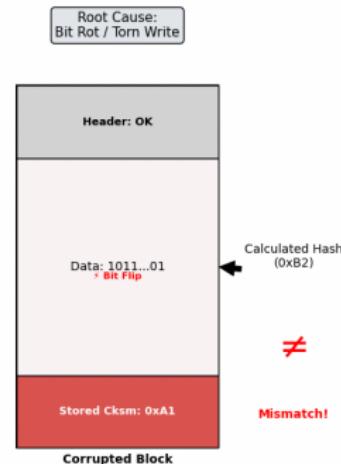
Data is read, hash is calculated, but

**Calculated Checksum  $\neq$  Stored Checksum.**

Type 1: Checksum Mismatch (CM)

## Root Causes:

- ▶ **Bit-level Corruption:** Media aging, head interference, bit rot.
- ▶ **Torn Writes:** Power failure during write (partial data written).
- ▶ **Misdirected Writes:** Overwriting good data at wrong location.



## Detection Mechanism:

- ▶ Detected during **ANY RAID Read** (User Read, Reconstruction, or Scrubbing).

# Type 2: Identity Discrepancies (IDs)

## Definition

The Checksum is **Valid** (Physical bits are intact), but the block's metadata (Inode/Offset) does not match the File System request.

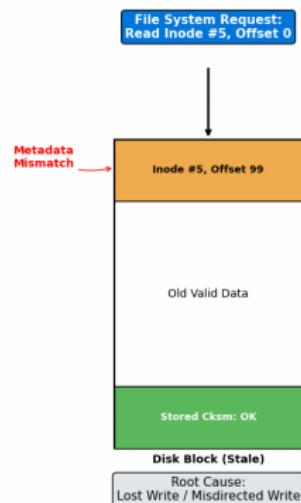
## Root Causes:

- ▶ **Lost Writes:** Disk reports "Success" but data was never written. System reads back *old, valid data*.
- ▶ **Misdirected Writes:** Data written to the wrong address.

## Detection Mechanism:

- ▶ **Only** detected during a **File System Read**.

Type 2: Identity Discrepancy (ID)



# Type 3: Parity Inconsistencies (PIs)

## Definition

All Data Block Checksums are **Valid**, but:

$$\text{Data}_1 \oplus \text{Data}_2 \oplus \dots \neq \text{Parity}$$

## Root Causes:

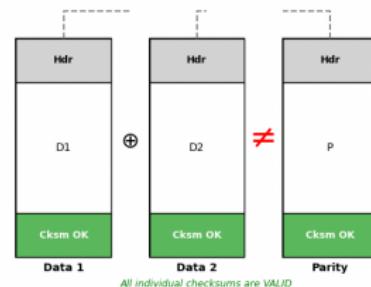
- ▶ **Memory Corruption:** Bit flips in RAID controller cache.
- ▶ **Software Bugs:** Errors in RAID algo.
- ▶ **Non-atomic Updates:** Power loss between Data write & Parity update.

## Detection Mechanism:

- ▶ **Only** detected during **Data Scrubbing** (Background verification).

Type 3: Parity Inconsistency (PI)

[RAID Logic / Memory]



Root Cause:  
Mem Corruption / SW Bug

## Challenge I: Prevalence Challenge

### Observation Data:

- ▶ Total Checksum Mismatches: **400,000+**
- ▶ SATA (Nearline) error rate is **10x** higher than FC (Enterprise).

### RAID Challenge 1: The Normalized Threat

- ▶ *RAID Assumption:* Data corruption is treated as a **rare and exceptional anomaly**.
- ▶ *Study Conclusion:* Data corruption is actually a **frequent and persistent operational threat**.
- ▶ **Impact:** RAID must evolve from an occasional recovery mechanism to an all-weather defense system.

## Challenge II: Reconstruction Hazard

### Observation Data:

- ▶ **8%** of errors were discovered during **RAID Reconstruction**.

### Logical Deduction:

- ▶ Reconstruction = System's most vulnerable moment (Redundancy Lost).
- ▶ Error during Reconstruction = **Double Failure**.

### RAID Challenge 2: Single Parity Failure

- ▶ RAID 5 cannot handle this 8% scenario.
- ▶ **Result:** Permanent Data Loss.

# Challenge III: Spatial Locality

Observation: Distribution of Errors

## Observation Data:

- ▶ Errors tend to cluster in **consecutive physical blocks**.
- ▶ Specific disk models fail at **specific block addresses**.

## RAID Challenge 3: Layout Vulnerability

- ▶ Traditional RAID uses **Aligned Stripes** (Same offset across disks).
- ▶ **Risk:** A firmware bug targeting a specific address can corrupt the *entire stripe* (Data + Parity) simultaneously.
- ▶ **Need:** Staggered Striping layouts.

## Challenge IV: System Correlation

Testing the Independence Assumption

### Hypothesis Testing:

- ▶  $H_0$ : Disk failures are independent.
- ▶ **Result:** Reject  $H_0$ . Failures are highly correlated.

### Evidence:

- ▶ Extreme case: **92 disks** in a single system failed simultaneously.

## RAID Challenge 4: The Independence Deficit

- ▶ **Cause:** Shared hardware (Controllers/Adapters) leads to common-mode failures.
- ▶ **Impact:** RAID groups must be distributed across physical failure domains.

# Lessons Learned I: Protection & Redundancy

## 1. The Value of Checksums

- ▶ **Observation:** Corruption is a reality, not a theory. We observed **400,000+** checksum mismatches. Up to **4%** of specific drive models developed errors within 17 months.
- ▶ **Conclusion:** The protection provided by **Checksums and Block Identity** is fully worth the extra storage space required.

## 2. The Reconstruction Hazard

- ▶ **Observation:** A significant portion (**8%**) of corruption is detected during RAID reconstruction.
- ▶ **Implication:**
  - ▶ Protection against **Double Disk Failures** (e.g., RAID 6) is necessary to prevent data loss.
  - ▶ **Aggressive Scrubbing** is crucial to detect errors *before* a rebuild is needed.

## Lessons Learned II: Policy & Layout

### 3. Enterprise Drive Replacement Policy

- ▶ **Observation:** Enterprise drives have a low probability of initial corruption, but once they fail, they are highly likely to fail again.
- ▶ **Strategy:** "One Strike Policy" It makes sense to **replace Enterprise drives immediately** upon the first detected corruption. Since the initial probability is low, the total replacement cost remains manageable.

### 4. Physical Layout Design

- ▶ **Observation:** Certain block numbers are more prone to corruption than others (due to firmware/hardware bugs affecting specific addresses).
- ▶ **Design Requirement: Staggered Striping** RAID designers should ensure stripe blocks are **NOT** stored at the same or nearby physical block numbers across disks.

# Thank You!